

## **Reinforcement Learning:**

**Reinforcement learning is a type of machine learning where an agent learns how to behave in an environment by performing actions and receiving feedback in the form of rewards or penalties.** The goal is for the agent to discover the best way to achieve a specific objective over time. The agent explores different actions, learns from the consequences, and adjusts its strategy to maximize cumulative rewards. It's like teaching a computer to make decisions through **trial and error**, with the ultimate aim of making better choices based on past experiences.

**Supervised learning is not always possible** due to challenges in obtaining labeled data, where each input has a corresponding correct output. Here are some simple examples:

### **1. Rare diseases diagnosis:**

- Example: Identifying extremely rare diseases may not have enough labeled cases, making it difficult to train a supervised model effectively.

### **2. Autonomous driving:**

- Example: Collecting labeled data for every possible driving scenario, including rare and dangerous situations, can be impractical and may not cover all possible conditions on the road.

### **3. Fraud detection:**

- Example: Fraudulent activities are relatively uncommon, and obtaining labeled data for various types of fraud instances may be limited, making it challenging to train a supervised model to accurately detect fraud.

### **4. Speech recognition for diverse accents:**

- Example: Training a supervised model for recognizing speech in various accents would require a vast amount of labeled data for each accent, which may not be readily available.

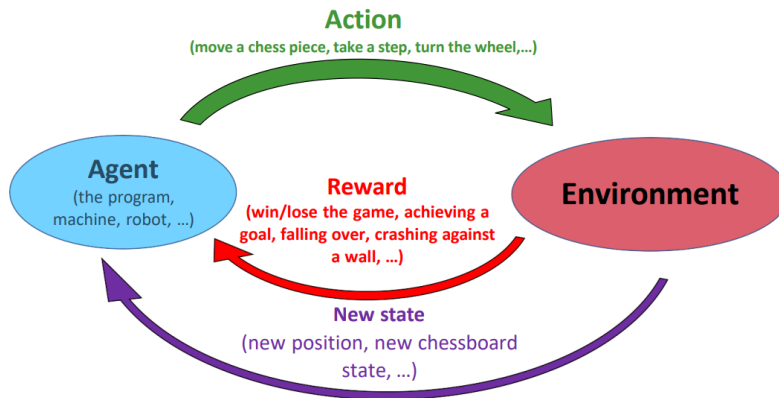
### **5. Predicting natural disasters:**

- Example: Anticipating rare and extreme events like earthquakes or tsunamis might lack sufficient labeled data due to the infrequency of such occurrences.

### **6. Sentiment analysis in new domains:**


- Example: If you want to perform sentiment analysis on user reviews for a completely new product or service, you may not have labeled data for that specific domain, making supervised learning challenging.

In these cases, it may be more practical to explore other machine learning approaches, such as unsupervised learning or reinforcement learning, which don't rely on having labeled examples for every possible scenario.



Flow Chart

### Explain the Cart-Pole Problem (Important For Final)



**Objective:** The goal is to balance the pole for as long as possible by making appropriate movements with the cart. The environment gives a reward for every time step the pole remains upright.

**State:** The agent (the algorithm or model) receives information about the current state of the system, including the **cart's position, velocity, the pole's angle, and its angular velocity.**

**Action:** The agent can take actions to move the cart left or right.

**Reward:** The agent learns through trial and error, receiving rewards(1) when it successfully keeps the pole upright and penalties when it fails.

#### **Challenges of reinforcement learning:**

- **Evaluative feedback:** Need trial and error to find the right action
- **Delayed feedback:** Actions may not lead to immediate reward
- **Non-stationarity:** Data distribution of visited states changes when the policy changes
- Fleeting nature of time and online data

### **What is Markov decision process?**

A Markov Decision Process (MDP) is a mathematical framework used to **model decision-making in situations where an agent interacts with an environment**. The key idea is the Markov property, which says that the future state of the system only depends on the **current state and the action taken**, not on the history of events leading up to it.

#### **Differences between Fully Observed MDP and Partially observed MDP:**

##### **Observability:**

Fully Observed MDP: Agent has direct access to the true and complete state of the environment.

Partially Observed MDP: Agent receives incomplete or noisy observations and maintains a belief over possible states.

##### **Decision-Making Challenge:**

Fully Observed MDP: Decision-making is relatively straightforward as the agent has perfect knowledge of the current state.

Partially Observed MDP: Decision-making is more challenging as the agent must deal with uncertainty and make decisions based on partial information.

##### **Representation:**

Fully Observed MDP: The state provides a complete and sufficient representation of the environment.

Partially Observed MDP: The agent needs to maintain a belief or probability distribution over possible states.

##### **Examples:**

Fully Observed MDP: **Robot navigation with accurate sensors.**

Partially Observed MDP: **Surveillance drone with limited visibility.**

An MDP is defined as a tuple  $(\mathcal{S}, \mathcal{A}, \mathcal{R}, \mathbb{T}, \gamma)$

$\mathcal{S}$  : Set of possible states

$\mathcal{A}$  : Set of possible actions

$\mathcal{R}(s, a, s')$  : Distribution of reward

$\mathbb{T}(s, a, s')$  : Transition probability distribution, also written as  $p(s'|s,a)$

$\gamma$  : Discount factor

### **Policy:**

A policy is a strategy or set of rules that guides the decision-making of an agent in different states. The goal of the agent is to find the best or optimal policy, which leads to the highest cumulative reward over time.

### ***Optimal Policy:***

An optimal policy is the best possible strategy an agent can follow to achieve the maximum expected cumulative reward in an MDP. It specifies the action the agent should take in each possible state to maximize its long-term rewards.

Finding the optimal policy is the central objective in many reinforcement learning problems.

A policy  $\pi^*$  is called **optimal** if it has maximal value for all states  $s \in \mathcal{S}$ :

$$V_{\pi^*}(s) = \max_{\pi} V_{\pi}(s).$$

**Discount Factor:** The discount factor, often denoted by the symbol ( $\gamma$ ), is a parameter used in reinforcement learning and Markov Decision Processes (MDPs). It influences the agent's decision-making by assigning different weights to immediate rewards compared to future rewards. The discount factor is a value between 0 and 1.

**Q: What is Additive discounted reward? How can you justify it?**

### **Additive Rewards in MDPs:**

In an MDP, the agent interacts with an environment by taking actions in different states. At each time step, the agent receives a reward based on the action taken and the resulting state transition. The total reward collected by the agent during its interaction with the environment is the sum of these individual rewards.

### **Discounted Rewards in MDPs:**

Discounted rewards involve applying a discount factor (usually denoted by  $\gamma$ , where  $0 \leq \gamma \leq 1$ ) to future rewards in order to account for the time preference of the agent. The discounted cumulative reward, or return, at a given time step  $t$  is calculated as the sum of  $\gamma^{(t-1)} * R_t$ , where  $R_t$  is the reward at time step  $t$ .

Additive discounted reward in an MDP is the sum of the discounted rewards over all time steps:

$$G_t = R_1 + \gamma R_2 + \gamma^2 R_3 + \dots + \gamma^{t-1} R_t$$

Discounting encourages the agent to **prioritize immediate rewards over delayed rewards**, which can be essential in situations where actions have long-term consequences and where there is uncertainty about the future.

Combining additive rewards with a discount factor in the context of an MDP helps in formulating a meaningful objective function for the agent, considering both the total reward and the temporal aspect of the decision-making process.

Certainly! Let's break down the statement and provide an easy example to illustrate the concept:

**Explanation:**

In decision-making, discounting involves giving less importance to rewards that are received in the future compared to those received immediately. This is particularly useful when the outcomes of actions have long-term effects and when there's uncertainty about what might happen in the future.

**Example:**

Imagine you have a choice between receiving \$100 today or receiving \$100 a year from now. If you choose to receive the money today, you are prioritizing immediate rewards over delayed rewards. This makes sense in situations where the future is uncertain, such as not being sure if you'll have the opportunity to receive the money a year later. Discounting reflects the idea that a reward received sooner is generally more valuable than the same reward received later.

**Explanation:**

In the context of Markov Decision Processes (MDP), which model decision-making over time, adding a discount factor to the total rewards helps create a meaningful objective function for the agent. This function considers both the overall rewards and the time aspect of decision-making.

**Example:**

Consider playing a video game where you earn points for completing tasks. The tasks completed earlier might contribute more to your total score than tasks completed later. This reflects the idea that the game rewards immediate achievements more than delayed ones. The discount factor in this case helps the agent (you, the player) to make decisions that balance the importance of both immediate and future rewards, leading to a more strategic and optimal decision-making process.

## Q-Learning:

**Exploration-exploitation** is a fundamental challenge in reinforcement learning, representing the trade-off between trying new actions to discover their effects (exploration) and choosing actions that are known to yield high rewards based on past experiences (exploitation). Striking the right balance between exploration and exploitation is crucial for an agent to learn an optimal policy in a reinforcement learning problem.

Here's a more detailed explanation:

### **1. Exploration:**

- Purpose: To discover the effects of different actions in various states.
- Strategy: The agent intentionally **selects** actions that **it hasn't tried much or doesn't have much information about**. This helps in learning more about the environment and improving the agent's understanding of the state-action space.
- Challenge: Too much exploration may **lead to suboptimal** immediate rewards since the agent is deliberately trying less-promising actions.

### **2. Exploitation:**

- Purpose: To choose actions that are expected to yield high rewards based on the agent's current knowledge.

- Strategy: The agent selects actions that it believes have been effective in the past or are likely to be the best based on its current understanding of the environment.

**Exploitation aims to maximize short-term rewards.**

- Challenge: Over-reliance on exploitation may cause the agent to **miss out on discovering better actions or adapting to changes** in the environment.

In practice, a common strategy for balancing exploration and exploitation is the  **$\epsilon$ -greedy** strategy:

- With probability  $\epsilon$  (exploration rate), the agent selects a random action.
- With probability  $(1 - \epsilon)$  (exploitation rate), the agent selects the action with the highest estimated value (Q-value) based on its current knowledge.

**This strategy allows the agent to explore new possibilities while still favoring actions that are deemed(মনে করা) promising based on its existing knowledge.**

The choice of the exploration rate ( $\epsilon$ ) is crucial, and it often decreases over time as the agent gains more experience. Initially, higher exploration rates help the agent explore the environment thoroughly, and as it learns, a shift towards more exploitation helps in maximizing cumulative rewards.

Algorithm:

## Q-learning + Epsilon-Greedy

**Q-learning algorithm:** Find for all  $s \in S$  and  $a \in A$  a function  $Q(s, a)$ , which gives a good approximation for  $Q^*(a, s)$ .

1. Start with random values for  $Q(s, a)$ . (e.g. all zero)
2. Choose a starting state  $s_0 \in S$ .
3. Look up the current best action in that state, i.e.  $a_0 = \operatorname{argmax}_{a \in A} Q(s_0, a)$  **or choose a random action  $a_0 \in A$  with probability  $\epsilon \in [0, 1]$  (Epsilon-Greedy Algorithm).**
4. Apply this action and get a new state  $s_1$  and reward  $r_0 = R(s_0, a_0, s_1)$ .
5. Update the value  $Q(s_0, a_0)$  as follows (**Bellman equation**)

$$Q(s_0, a_0) = (1 - \alpha)Q(s_0, a_0) + \alpha \left( r_0 + \gamma \max_{a \in S} Q(s_1, a) \right).$$

Here  $\alpha \in [0, 1]$  is the **learning rate**.

6. If  $s_1$  is not a terminal state repeat with step 3.

The parameter  $\epsilon$  controls the exploration rate, and it is a hyperparameter that can be tuned based on the learning task and environment.

**Q: Discuss the role and intuition of each component of the update (state and action) equation.**

In the Q-learning algorithm, the update equation is a crucial component that guides how the Q-values, denoted as  $Q(s_0, a_0)$ , are adjusted based on the observed rewards and the maximum expected future rewards. The update equation has several components, each serving a specific role in shaping the learning process. Let's break down the update equation:

$$Q(s_0, a_0) = (1 - \alpha)Q(s_0, a_0) + \alpha \left( r_0 + \gamma \max_{a \in S} Q(s_1, a) \right).$$

### 1. $Q(s_0, a_0)$ :

- Role: Represents the current estimate of the Q-value for taking action  $a_0$  in state  $s_0$ .
- Intuition: This is the value being updated. It reflects the agent's current belief about the cumulative future rewards associated with taking action  $a_0$  in state  $s_0$ .

### 2. $\alpha$ :

- Role: Learning rate, a hyperparameter between 0 and 1 that controls the **step size of the update**.

-Intuition: The learning rate determines how much the new information should influence the current estimate. A **higher learning rate means the new information has a larger impact on the Q-value.**

### 3. $r_0$

- Role: Immediate reward observed after taking action  $a_0$  in state  $s_0$ .

-Intuition: Represents the immediate reinforcement obtained by the agent. The agent uses this reward to update its Q-value, incorporating the knowledge of the immediate consequences of the action.

### 4. gamma:

- Role: Discount factor, a hyperparameter between 0 and 1 that **discounts the impact of future rewards.**

**The discount factor is like a weighing scale for rewards over time. A value of 0 means the agent only cares about immediate rewards, while a value of 1 means the agent equally values both immediate and future rewards.**

-Intuition: Encourages the agent to consider not only immediate rewards but also future rewards. A high gamma values long-term rewards more, and a low gamma values short-term rewards more.

### 5. $\max Q(s_1, a)$ :

- Role: Represents the maximum Q-value for the next state  $s_1$  **over all possible actions  $a$ .**

- Intuition: This term captures the expected future rewards. The agent estimates the maximum potential cumulative **ক্রমবর্ধমান** reward it can obtain from the next state  $s_1$  by selecting the action that maximizes the Q-value.

## Clustering:

**Advantages and disadvantages of min(Single Link) and max(Complete Link) of inter-cluster distance:**

### Single Linkage:

#### **Advantages:**

**1.Sensitivity to local structures:** Single linkage tends to form clusters based on local structures and is more likely to capture elongated and irregularly shaped clusters.

**2.Early merging of nearby points:** It is computationally efficient and often results in early merging of nearby data points, which can be advantageous in certain scenarios.



### ***Disadvantages:***

- 1. Chain effect:** Single linkage is sensitive to outliers and noise, leading to a "chain effect" where outliers can join clusters in a chain-like manner, connecting distant points.
- 2.Susceptible to chaining:** It can create long chains of data points even if they are not part of the same cluster, making it susceptible to the chaining phenomenon.

### **Complete Linkage:**

#### ***Advantages:***

- 1.Robust to outliers:** Complete linkage is less sensitive to outliers compared to single linkage, as it considers the maximum distance between points from different clusters.
- 2.Compact and spherical clusters:** It tends to form more compact and spherical clusters, which can be advantageous in scenarios where clusters have a more globular shape.

#### ***Disadvantages:***

- 1.Merge bias:** Complete linkage tends to have a merge bias towards globular clusters and may not perform well with clusters that have irregular shapes or elongated structures.
- 2.Tendency to produce unbalanced clusters:** It may lead to unbalanced clusters, where one cluster is substantially larger than the others, particularly if there are outliers or noise.

c) Consider the dataset given below. Cluster the records using agglomerative hierarchical clustering.

ID	A1	A2	A3
1	0	0	0
2	1	1	0
3	1	1	1
4	1	1	0

Define the distance metric using Manhattan distance between clusters. Show the intermediate result using single link and the final result in dendrogram.

Integer-43

3c) using manhattan;

$$d(1,2) = d(2,1) = |0-1| + |0-1| + |0-0| = 2$$

$$d(1,3) = d(3,1) = |0-1| + |0-1| + |0-1| = 3$$

$$d(1,4) = d(4,1) = |0-1| + |0-1| + |0-0| = 2$$

$$d(2,3) = d(3,2) = |1-1| + |1-1| + |0-1| = 1$$


$$d(2,4) = d(4,2) = |1-1| + |1-1| + |0-0| = 0$$

$$d(3,4) = d(4,3) = |1-1| + |1-1| + |1-0| = 1$$

↓

Dist	1	2	3	4
1	0	2	3	2
2	2	0	1	0
3	3	1	0	1
4	2	0	1	0

Dendrogram:



using Single linkage;

↓

min dist	1	(2,4)	3
1	0	2	3
(2,4)	2	0	1
3	3	1	0

↓

mindist	1	((2,4),3)
1	0	2
((2,4),3)	2	0

∴ The minimum distance is 2.

## PCA and LDA:

### Scenarios in which PCA works better:

**Exploratory Data Analysis:** PCA is well-suited for **exploratory data** analysis when the primary goal is to understand the overall structure and relationships within the data. It helps in identifying patterns and reducing the dimensionality of the data.

**Unsupervised Learning:** PCA is an unsupervised technique, making it suitable for situations where class labels or groupings are not available or not relevant. It is often used for **feature extraction** in tasks **such as image compression or denoising**.

**Dimensionality Reduction for Efficiency:** When computational efficiency is a concern, PCA is a good choice, as it is computationally efficient and can handle large datasets.

### Scenarios in which LDA works better:

**Classification Tasks:** LDA is specifically designed for classification problems. When the goal is to maximize the separation between classes, LDA often outperforms PCA because it takes class information into account during dimensionality reduction.

**Supervised Learning with Labeled Data:** LDA requires labeled data, and it is most effective when there is a clear distinction between classes. It is commonly used in **face recognition, speech recognition, and other classification tasks**.

**Preserving Discriminant(বৈষম্যমূলক)Information:** If the goal is to preserve information that helps discriminate between classes, LDA is the preferred choice. It **optimizes for both variance** within classes and separation between classes.

### Limitations of LDA:

#### **1. Limited Number of Feature Projections:**

**Explanation:** LDA can generate at most  $C-1$  feature projections, where  $C$  is the number of classes. This means that if the classification error suggests the need for more features, additional methods must be used to provide those extra features.

**Example:** If you have a dataset with three classes ( $C=3$ ), LDA can generate at most two feature projections.

## ***2. Parametric Assumption - Unimodal Gaussian Likelihoods:***

**Explanation:** LDA assumes that the data within each class follows a unimodal Gaussian (normal) distribution. If the actual distributions significantly deviate বিচ্যুত from this assumption and are non-Gaussian, LDA may not effectively preserve the complex structure of the data needed for classification.

**Example:** If the data within a class has a distribution that is not bell-shaped like a normal distribution, LDA might not perform optimally.

## ***3. Failure when Discriminatory Information is in Variance, Not Mean:***

**Explanation:** LDA assumes that the discriminatory information is primarily in the mean of the data. If the crucial information for classification is in the variance (spread) of the data rather than the mean, LDA may fail to capture this essential information.

**Example:** In situations where the differences between classes are better represented by the spread or variability of the data rather than the average values, LDA may not be the most suitable method.

**PCA MATH:**

- e) The dataset shown in Table 1, reports the price and weight of three products. Show the step-by-step calculation to derive the feature vector by applying principal component analysis (PCA).

Table 1

Product	Price	Weight
1	4	7
2	5	6
3	3	1

Integer-43

3(e)

Data:

Product	Price	Weight
1	4	7
2	5	6
3	3	1

Step-1: Get some data

$$\begin{aligned}
 X &= 12 & Y &= 14 \\
 \bar{X} &= 4 & \bar{Y} &= 4.67 \\
 n &= 3
 \end{aligned}$$

Step-2: Subtract the mean

	$X_i - \bar{X}$	$Y_i - \bar{Y}$
1	1	2.33
2	2	1.33
3	0	-3.67

Step-3: Calculate cov

$$\text{cov} = \begin{pmatrix} 2.5 & 6.25225 \\ 6.25225 & 10.33335 \end{pmatrix} = A$$

Step-4: Calculate eigenvalues and eigenvectors

$$\begin{aligned}
 A - \lambda I &= \begin{pmatrix} 2.5 & 6.25225 \\ 6.25225 & 10.33335 \end{pmatrix} - \lambda \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix} \\
 &= \begin{pmatrix} 2.5 - \lambda & 6.25225 \\ 6.25225 & 10.33335 - \lambda \end{pmatrix} = 0
 \end{aligned}$$

$$\Rightarrow (2.5 - \lambda)(10.33335 - \lambda) - 39.09063006 = 0$$

$$\Rightarrow 25.833375 - 2.5\lambda - 10.33335\lambda + \lambda^2 - 39.09063006 = 0$$

$$\Rightarrow \lambda^2 - 12.83335\lambda - 13.25725506 = 0$$

$$\Rightarrow \lambda_1 = 13.79440996$$

$$\lambda_2 = -0.9610599583$$

$$\therefore \text{eigenvalues} = \begin{pmatrix} 13.79440996 \\ -0.9610599583 \end{pmatrix}$$

Step-5: choosing Components and forming feature vectors

For  $\lambda_1 = 13.79440996$

$$\begin{pmatrix} 2.5 - 13.79440996 & 6.25225 \\ 6.25225 & 10.33335 - 13.79440996 \end{pmatrix} \begin{bmatrix} x_1 \\ y_1 \end{bmatrix} = 0$$

$$= \begin{pmatrix} -11.29440996 & 6.25225 \\ 6.25225 & -3.46105996 \end{pmatrix} \begin{bmatrix} x_1 \\ y_1 \end{bmatrix} = 0$$

Now,

$$-11.29440996 x_1 + 6.25225 y_1 = 0 \quad \text{--- (i)}$$

$$6.25225 x_1 - 3.46105996 y_1 = 0 \quad \text{--- (ii)}$$

$$\text{(i)} \Rightarrow x_1 = 0.5535 y_1$$

$$\text{(ii)} \Rightarrow x_1 = 0.5535 y_1$$

$$e_1 \sim \begin{bmatrix} 0.5535 \\ A \\ \frac{1}{A} \end{bmatrix}$$

$$\therefore e_1 \sim \begin{bmatrix} 0.4842 \\ 0.8749 \end{bmatrix}$$

$$A = \sqrt{(0.5535)^2 + (1)^2} = 1.1429$$



For  $\lambda_2 = -0.9610599583$

$$\begin{pmatrix} 25 + 0.9610599583 & 6.25225 \\ 6.25225 & 10.33335 + 0.9610599583 \end{pmatrix} \begin{bmatrix} x_1 \\ y_1 \end{bmatrix} = 0$$

$$\begin{pmatrix} 3.461059 & 6.25225 \\ 6.25225 & 11.29440 \end{pmatrix} \begin{bmatrix} x_1 \\ y_1 \end{bmatrix} = 0$$

Now,

$$3.461059 x_1 + 6.25225 y_1 = 0 \quad \text{--- (II)}$$

$$6.25225 x_1 + 11.29440 y_1 = 0 \quad \text{--- (IV)}$$

$$\text{(II)} \Rightarrow x_1 = -1.8064 y_1$$

$$\text{(IV)} \Rightarrow y_1 = -1.8064 x_1$$

$$e_2 \sim \begin{bmatrix} \frac{1.8064}{A} \\ \frac{1}{A} \end{bmatrix}$$

$$A = \sqrt{(1.8064)^2 + (1)^2} = 2.06$$

$$\therefore e_2 \sim \begin{bmatrix} 0.8768 \\ 0.4851 \end{bmatrix}$$

$\therefore$  The feature vector is  $\begin{bmatrix} 0.4842 \\ 0.8749 \end{bmatrix}$  for  $\lambda_1 = 13.7944$

## LDA MATH:

### Question 4. [Marks: 14]

✓ Compute the Linear Discriminant projection for the following two dimensional datasets. [10+4]

$$\omega_1: X_1 = (x_1, x_2) = \{(5,1), (1,3), (4,3)\}$$

$$\omega_2: X_2 = (x_1, x_2) = \{(8,9), (8,6), (10,5)\}$$

PCA & LDA

Also, is it possible to combine both PCA and LDA together to get a better result? Justify your answer with appropriate logic.



## LDA

### Enigma-41

4

$$\omega_1: X_1 = (x_1, x_2) = \{(5, 1), (1, 3), (4, 3)\}$$

$$\omega_2: X_2 = (x_1, x_2) = \{(8, 9), (8, 6), (10, 5)\}$$

For class - 1 ( $\omega_1$ ):

$x_1$	$x_2$
5	1
1	3
4	3

$$\bar{x}_1 = 3.33 \quad ; n=3$$

$$\bar{x}_2 = 2.33$$

$$\therefore \mu_1 = \begin{pmatrix} 3.33 \\ 2.33 \end{pmatrix}$$

cov of  $\omega_1$ :

$$S_1 = \text{Cov}(X_1) = \begin{pmatrix} 4.33335 & -1.66665 \\ -1.66665 & 1.33335 \end{pmatrix}$$

cov of  $\omega_2$ :

$$S_2 = \text{Cov}(X_2) = \begin{pmatrix} 1.33335 & -1.66665 \\ -1.66665 & 4.33335 \end{pmatrix}$$

Within Scatter class:

$$S_w = S_1 + S_2 = \begin{pmatrix} 5.6667 & -3.3333 \\ -3.3333 & 5.6667 \end{pmatrix}$$

For class - 2 ( $\omega_2$ ):

$x_1$	$x_2$
8	9
8	6
10	5

$$\bar{x}_1 = 8.67 \quad ; n=3$$

$$\bar{x}_2 = 6.67$$

$$\therefore \mu_2 = \begin{pmatrix} 8.67 \\ 6.67 \end{pmatrix}$$

Between class scatter matrix:

$$\begin{aligned}
 S_B &= (\mu_1 - \mu_2)(\mu_1 - \mu_2)^T \\
 &= \begin{bmatrix} 3.33 \\ 2.33 \end{bmatrix} - \begin{bmatrix} 8.67 \\ 6.67 \end{bmatrix} \begin{bmatrix} 3.33 \\ 2.33 \end{bmatrix} - \begin{bmatrix} 8.67 \\ 6.67 \end{bmatrix}^T \\
 &= \begin{bmatrix} -5.34 \\ -4.34 \end{bmatrix} \begin{bmatrix} -5.34 & -4.34 \end{bmatrix} \\
 &= \begin{pmatrix} 28.516 & 23.176 \\ 23.176 & 18.836 \end{pmatrix}
 \end{aligned}$$

Now,

$$S_B^{-1} S_B \omega = \lambda \omega$$

$$\therefore |S_B^{-1} S_B - \lambda I| = 0$$

$$\Rightarrow \begin{vmatrix} \begin{pmatrix} 5.6667 & -3.3333 \\ -3.3333 & 5.6667 \end{pmatrix}^{-1} \begin{pmatrix} 28.516 & 23.176 \\ 23.176 & 18.836 \end{pmatrix} - \lambda \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix} & \\ & \end{vmatrix} = 0$$

$$\Rightarrow \begin{vmatrix} \begin{pmatrix} 0.26988 & 0.15873 \\ 0.15873 & 0.26984 \end{pmatrix} \begin{pmatrix} 28.516 & 23.176 \\ 23.176 & 18.836 \end{pmatrix} - \begin{pmatrix} \lambda & 0 \\ 0 & \lambda \end{pmatrix} & \\ & \end{vmatrix} = 0$$

$$\Rightarrow \begin{vmatrix} 11.3746 - \lambda & 9.2445 \\ 10.78 & 8.7614 - \lambda \end{vmatrix} = 0$$

$$\Rightarrow (11.3746 - \lambda)(8.7614 - \lambda) - 99.66 = 0$$

$$\Rightarrow 99.66 - 11.3746\lambda - 8.7614\lambda + \lambda^2 - 99.66 = 0$$

$$\Rightarrow \lambda^2 - 20.136\lambda = 0$$

$$\Rightarrow \lambda(\lambda - 20.136) = 0$$

$$\therefore \lambda_1 = 0 \quad \text{and} \quad \lambda_2 = 20.136$$

For  $\lambda_2 = 20.136$ ;

$$\begin{pmatrix} 11.3746 - 20.136 & 9.2445 \\ 10.78 & 8.7614 - 20.136 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = 0$$

$$-8.7614x_1 + 9.2445x_2 = 0 \quad \dots \textcircled{I}$$

$$10.78x_1 - 11.3746x_2 = 0 \quad \dots \textcircled{II}$$

$$\textcircled{I} \Rightarrow x_1 = 1.055x_2$$

$$\textcircled{II} \Rightarrow x_2 = 1.055x_1$$

$$\text{So, } e \sim \begin{bmatrix} \frac{1.055}{A} \\ \frac{1}{A} \end{bmatrix}$$

$$\text{Here, } A = \sqrt{(1.055)^2 + 1} \\ = 1.4537$$

$$\therefore e \sim \begin{bmatrix} 0.726 \\ 0.687 \end{bmatrix} \quad \text{for max } \lambda.$$

For  $\lambda_1 = 0$ ;

$$11.3746x_1 + 9.2445x_2 = 0 \quad \dots \textcircled{III}$$

$$10.78x_1 + 8.7614x_2 = 0 \quad \dots \textcircled{IV}$$

$$\textcircled{III} \Rightarrow x_1 = -0.813x_2$$

$$\textcircled{IV} \Rightarrow x_1 = -0.813x_2$$

$$\text{So, } e \sim \begin{bmatrix} \frac{-0.813}{A} \\ \frac{1}{A} \end{bmatrix}$$

$$A = \sqrt{(-0.813)^2 + 1} \\ = 1.289$$

$$e \sim \begin{bmatrix} -0.631 \\ 0.776 \end{bmatrix}$$

Yes, it is possible to combine PCA and LDA to get a better result in certain scenarios.

- PCA can be applied initially to reduce dimensionality and capture the overall variance in the data.
- The reduced dataset from PCA can then be used as input for LDA, which focuses on maximizing class separability in the lower-dimensional space.
- This combination leverages both techniques for noise reduction (PCA) and discriminative projection (LDA).

Justification:

PCA can help eliminate noise and capture the most significant variability in the data, providing a cleaner input for LDA. LDA, in turn, can then focus on the discriminating information within the reduced feature space, potentially improving classification performance.

In summary, the combination of PCA and LDA can be beneficial in situations where both noise reduction and class separability are important considerations.

## Boosting:

Bias-Variance Tradeoff: [Bias-Variance Trade Off - Machine Learning - GeeksforGeeks](#)

Boosting

Adaboost Algo:

Explain the algorithm line by line

① For  $i=1$  to  $N$  <sup>size of dataset</sup> Initialize the data weight  $w_1^{(i)} = \frac{1}{N}$

② For  $t=1$  to  $T$

Ⓐ Find a classifier  $h_t(x)$  by minimizing the weighted error  $J_t$

$$J_t = \sum_{i=1}^N w_t^{(i)} \times \underbrace{I(y^{(i)} \neq h_t(x^{(i)}))}_{\text{misclassification}}$$

multiplying the weights of them

Ⓑ Find the weighted error of  $h_t(x)$ :

$$\epsilon_t = \frac{\sum_{i=1}^N w_t^{(i)} \times I(y^{(i)} \neq h_t(x^{(i)}))}{\sum_{i=1}^N w_t^{(i)}}$$

and a new component is assigned weight based on its error:

$$\alpha_t = \ln((1 - \epsilon_t) / \epsilon_t)$$

We will assign some points to the classifier. We know that, in between the "good classifier that misclassify" and "bad classifier that misclassify", the first one is bad, so we need to give priority to that one. After merging all classifier voting is done and the one with higher  $\alpha$  value will get higher priority.

For weak classifier;  $\epsilon_t = 0.5$

$$\ln(1) = 0$$

So,  $\alpha = 0$ ; voting power is reduced for weak classifier.

For strong classifier;  $\alpha$  is higher.



c) The normalized weights are updated:

$$w_{t+1}^{(i)} = \underbrace{w_t^{(i)}}_{\text{old weight}} \underbrace{e^{\alpha_t I(y^{(i)} \neq t_t(x^{(i)}))}}_{\text{misclassified data}}$$

$e$  was taken to reduce error exponentially.


Here, For weak classifier, the weight will stay same.  
For strong " " " will be updated.

3) Combined classifier  $\hat{y} = \text{sign}(H_T(x))$  where  $H_T(x) = \sum_{t=1}^m \alpha_t h_t(x)$

## Integer-43:

e) Write the differences between bagging and boosting methods of ensemble learning. Define weak and strong classifiers. How does the AdaBoost algorithm assign weights to different classifiers so that combined classification error is minimized?

- Bagging is a method of merging the same type of predictions. Boosting is a method of merging different types of predictions.
- Bagging decreases variance, not bias, and solves over-fitting issues in a model. Boosting decreases bias, not variance.
- In Bagging, each model receives an equal weight. In Boosting, models are weighed based on their performance.
- Models are built independently in Bagging. New models are affected by a previously built model's performance in Boosting.
- In Bagging, training data subsets are drawn randomly with a replacement for the training dataset. In Boosting, every new subset comprises the elements that were misclassified by previous models.
- Bagging is usually applied where the classifier is unstable and has a high variance. Boosting is usually applied where the classifier is stable and simple and has high bias.



### Weak Classifier:

A weak classifier is a model that performs slightly better than random chance but is not highly accurate on its own. It may make errors and is often considered a simple or basic model. Weak classifiers are typically used in ensemble methods like boosting, where multiple weak classifiers are combined to create a strong classifier.

### Strong Classifier:

A strong classifier is a model that is highly accurate and performs well on its own. It has the ability to capture complex patterns and make accurate predictions. Strong classifiers are often the result of combining multiple weak classifiers or through the use of sophisticated algorithms.

The algorithm starts by assigning equal weights to all training instances. In each iteration, a weak learner is trained on the weighted training data, and its classification error is computed. The weight of each instance is then updated based on whether it was classified correctly or not. The weight of misclassified instances is increased, while the weight of correctly classified instances is decreased. This process is repeated for a fixed number of iterations or until the desired accuracy is achieved. The final classifier is a weighted combination of all the weak learners, where the weights are proportional to their accuracy.

## Template Matching:

- 6) Explain Bellman's optimality principle. Find the edit distance between the word "pattern" and its misspelled version "petern". What is the sequence of edit operations to achieve them?

Integer-43  
3(b)

Bellman's optimality principle:

$$i_0, j_0 \xrightarrow{\text{opt}} i_f, j_f \text{ can be obtained as}$$

$$i_0, j_0 \xrightarrow{\text{opt}} i, j \oplus i, j \xrightarrow{\text{opt}} i_f, j_f$$

The overall optimal path from  $i_0, j_0$  to  $i_f, j_f$  through  $i, j$  is the concatenation of the optimal paths from  $i_0, j_0$  to  $i, j$  and  $(i, j)$  to  $i_f, j_f$ .

		p	a	t	t	e	r	n
	0	1	2	3	4	5	6	7
p	1	0	1	2	3	4	5	6
e	2	1	1	2	3	3	4	5
t	3	2	2	1	2	3	4	5
e	4	3	3	2	2	2	3	4
r	5	4	4	3	3	3	2	3
n	6	5	5	4	4	4	3	2

Edit operation: Substitute (a-e), insert (t).

Optimal Path distance: 2



## SVM

- For simple, linear problems: A **Perceptron** is usually sufficient.
- For complex problems where relationships between data points can be modeled with a non-linear boundary and you have a lot of data: An **MLP** is often the best choice.
- For problems with clear margins between classes, whether linear or not, especially when the dataset is not too large or when the dimensionality is high: An **SVM** can be the most effective.
- **Naive Bayes** can be the algorithm of choice when the dataset is massive, and the speed is a concern, even if it may not provide the most accurate results. It can also perform well as a baseline in natural language processing tasks or in scenarios where interpretability is more important than achieving the last percentage of accuracy. Naive Bayes can be preferable when the dataset has missing values or when quick prototyping is needed.

What is the motivation of the support vector machine (SVM) based learning over a decision boundary (e.g., perceptron) based learning? Derive the equation that you want to minimize to get the solution of the SVM.

The motivation behind Support Vector Machines (SVM) is to find the optimal hyperplane that maximally separates the classes in the feature space. Unlike the perceptron, which merely finds any hyperplane that separates the classes, the SVM is concerned with the margin—the distance between the hyperplane and the nearest data point from either set. The larger the margin, the lower the generalization error of the classifier.

The SVM operates on the principle of margin maximization. Given a set of training examples, each marked as belonging to one of two categories, an SVM training algorithm builds a model that assigns new examples into one category or the other. It's a non-probabilistic binary linear classifier, although methods such as the kernel trick can enable SVM to perform non-linear classification.

## Bayesian

### Question 3. [Marks: 14]

- a) How do the generative learning algorithms differ from the discriminative learning algorithms? [2+2+2]  
What is the strong assumption of Naïve Bayes? Why Laplace smoothing factor is necessary for Naïve Bayes? ch2

---

#### Solution:

##### Generative Models:

Generative models are machine learning models that learn to generate new data samples similar to the training data they were trained on. They capture the underlying distribution of the data and can produce novel instances. Generative models find applications in image synthesis, data augmentation, and generating realistic content like images, music, and text.

##### Discriminative Models:

The discriminative model refers to a class of models used in Statistical Classification, mainly used for supervised machine learning. These types of models are also known as conditional models since they learn the boundaries between classes or labels in a dataset. Discriminative models focus on modeling the decision boundary between classes in a classification problem. The goal is to learn a function that maps inputs to binary outputs, indicating the class label of the input. Maximum likelihood estimation is often used to estimate the parameters of the discriminative model, such as the coefficients of a logistic regression model or the weights of a neural network.

The strong assumption of Naïve Bayes is that the features used in the model are conditionally independent given the class label. In other words, it assumes that the presence or value of one feature does not depend on the presence or value of any other feature, given the class label.

Laplace smoothing involves adding a small constant (usually 1). This "smoothing" ensures that no probability is exactly zero.

Laplace smoothing factor:

unsmoothed

$$p(X_1 = T \mid Y = \text{spam}) = \frac{\text{freq. of } T \text{ in spam data}}{\text{total \# of spam instances}}$$

smoothed

$$p(X_1 = T \mid Y = \text{spam}) = \frac{\text{freq. of } T \text{ in spam} + 1}{\text{total \# of spam instances} + v}$$

## Differences between MLE MAP bayesian estimated method

### Maximum Likelihood Estimation (MLE):

- Focuses solely on the likelihood of observing the given data under different parameter values.
- Does not incorporate prior knowledge about the parameters.
- Look for the parameter values that maximize the likelihood function, i.e., the values that make the observed data most probable.
- Assumes all parameter values are equally likely before observing any data (non-informative prior).

### Maximum A Posteriori (MAP):

- Combines the likelihood of the observed data with a prior distribution that encodes existing knowledge about the parameter values.
- Searches for parameter values that maximize the product of the likelihood and the prior probability (the posterior distribution).
- Can be biased by the prior, especially with small sample sizes; with a large number of samples, the influence of the prior diminishes, and MAP converges to MLE.

### Bayesian Estimation:

- A more general approach that also uses prior information.
- Rather than finding a single estimate, Bayesian estimation calculates the entire posterior distribution of the parameters, considering both the data (likelihood) and the prior knowledge.
- Allows for uncertainty in parameter estimates to be expressed and updated as new data is observed.

## Q:

What is Maximum a Posteriori (MAP) estimation in statistical inference and how does it differ from Maximum Likelihood Estimation (MLE)? [4]

### MAP Estimation:

MAP estimation incorporates prior knowledge about the parameters through a prior distribution. When estimating a parameter  $\theta$ , MAP takes into account the likelihood of the observed data given  $\theta$  as well as the prior probability of  $\theta$ . It selects the value of  $\theta$  that maximizes the product of these two quantities, which is equivalent to maximizing the posterior probability  $P(\theta|\text{data})$ .

Mathematically, the MAP estimate is given by:

$$\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} P(\theta|\text{data}) = \arg \max_{\theta} \frac{P(\text{data}|\theta)P(\theta)}{P(\text{data})}$$

Since  $P(\text{data})$  is constant with respect to  $\theta$ , it simplifies to:

$$\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} P(\text{data}|\theta)P(\theta)$$

### MLE Estimation:

MLE focuses solely on the likelihood of the observed data given the parameters, without considering any prior distribution. It selects the value of  $\theta$  that makes the observed data most probable. In other words, it maximizes the likelihood function  $P(\text{data}|\theta)$ .

The MLE estimate is given by:

$$\hat{\theta}_{\text{MLE}} = \arg \max_{\theta} P(\text{data}|\theta)$$

### Differences:

- **Incorporation of Prior Knowledge:** MAP uses prior knowledge about the parameters (prior distribution), while MLE does not.
- **Objective Function:** MAP maximizes the posterior distribution  $P(\theta|\text{data})$ , whereas MLE maximizes the likelihood function  $P(\text{data}|\theta)$ .
- **Estimate Interpretation:** MAP's estimate can be seen as a compromise between the likelihood of the data and the prior belief, while MLE's estimate is based purely on the data's likelihood.
- **Behavior with Data Size:** As the size of the dataset increases, the influence of the prior in MAP diminishes, and the MAP estimate converges to the MLE estimate, assuming the prior is not dogmatic.

## When bayesian distance classifier becomes minimum distance classifier

The Bayesian distance classifier becomes equivalent to a minimum distance classifier when all the class-conditional densities are identical and the cost of misclassification is the same for all classes. This simplification happens because when class-conditional densities are identical, the Bayesian decision rule simplifies to choosing the class with the nearest mean to the given data point, since the complex part of the decision which involves comparing likelihoods becomes unnecessary.

1. The Bayesian classifier makes a decision by computing posterior probabilities for each class given the data point and chooses the class with the highest posterior probability. This posterior probability is calculated using Bayes' theorem, which incorporates the likelihood of the data point given the class, the prior probability of the class, and the evidence (or marginal probability of the data point).
2. The likelihood term in Bayes' theorem is determined by the class-conditional density  $p(x|\omega_i)$ . If all classes have the same class-conditional density (for example, if the variance is the same across all classes in a Gaussian distribution), this term does not affect the comparison between classes.
3. When the cost of misclassification is the same for all classes, the decision rule becomes independent of the misclassification costs, and we can ignore them in the decision process.
4. Under these conditions, the posterior probability is proportional to the prior probability of the class. If all classes have the same prior probability, the classifier then simply measures the distance from the data point to the mean of each class and picks the class to which the data point is closest.

Cluster

**Q: Why normalization is bad for data when the variance is present due to cluster ?**

Ans: Normalization is a technique used to adjust the scale of different features in data so that they can be compared on common grounds. It typically involves scaling the features so that they have a specific range, like 0 to 1, or so that they have a mean of 0 and a standard deviation of 1.

However, **normalization might not be suitable in cases where the variance is important to preserve**, such as when the data contains inherent clusters. The reason is that normalization can sometimes distort the relative distances between points, which are critical for clustering algorithms to identify distinct groups.

Let's consider a simple example:

Suppose we have two clusters of data points:

Cluster 1: Points are close to each other, ranging from 90 to 110.

Cluster 2: Points are also close to each other, but range from 190 to 210.

The variance within each cluster is small, but the variance between clusters is large because the clusters are far apart.

If we apply min-max normalization to these points, scaling them to a range of 0 to 1, both clusters will be squished into a small range near each other. The normalized values will be close to 0 for Cluster 1 and close to 1 for Cluster 2. This might lead us to lose the understanding that

there were two distinct clusters quite far from each other, as the normalization process would make them appear as though they are similar.

In clustering problems, it's important that algorithms like K-means or hierarchical clustering can detect and maintain the true distance and variance between points. If normalization overly simplified these distances, the clusters that the algorithms find might not represent the true structure of the data. That's why in such cases, it's important to carefully consider whether to normalize the data or to use it in its original scale.

## Basic differences between partitioning and hierarchical clustering

Aspect	Hierarchical Clustering	Partitional Clustering
Cluster Structure	Provides a hierarchical structure of clusters	Provides a flat partitioning of the data into clusters
Number of Clusters	No need to pre-specify the number of clusters	Requires pre-specification of the number of clusters (K)
Computational Complexity	Computationally expensive for large datasets	Generally more scalable for large datasets
Handling Noisy Data	Maybe less robust to noisy or outlier data points	Can be more robust through iterative optimization
Flexibility	Enables exploration at different levels of granularity	Provides a fixed partitioning with less granularity
Initialization	Initialization is not required	Requires initialization of cluster centroids or parameters
Algorithm Examples	Agglomerative clustering, divisive clustering	K-means clustering, Expectation-Maximization (EM) algorithm

**The Design Cycle in machine learning and pattern recognition is a structured approach to creating a model that can accurately predict or classify data. It involves several steps, which can be iterated upon to refine the model. Here's a brief explanation of each step based on the slides you provided:**

#### **Data Collection:**

This step involves gathering the data that will be used to train and test the model. The data must be representative of the problem space to ensure that the model can generalize well to new, unseen data.

#### **Feature Selection:**

Once data is collected, the next step is selecting the features that will be used to train the model. This includes choosing the most informative attributes of the data that contribute to the predictive power of the model. The features should be discriminative, invariant to irrelevant transformations, and robust to noise and occlusion.

#### **Model Selection:**

After features are selected, a model must be chosen. This could be a parametric model, a non-parametric model, or a combination of different types. The choice depends on the problem domain, the amount of data, the complexity of the problem, and computational resources.

#### **Training:**

With the data and model selected, the next step is to train the model. This can involve supervised learning (where the model is trained with labeled data), unsupervised learning (where the model looks for patterns in the data without labels), or reinforcement learning (where the model learns by receiving feedback from interactions with the environment).

#### **Evaluation:**

The trained model needs to be evaluated to ensure it performs well. This can involve estimating performance with training samples (using metrics like accuracy, precision, recall) and predicting performance with future data (through methods like cross-validation to prevent overfitting and ensure generalization).

#### **Challenges:**

There are several challenges in the design cycle that can affect the performance of the model. These include occlusion (objects being blocked from view), scale (size variations of the objects), deformation (changes in shape), viewpoint variation (changes in the angle of observation), and illumination (variations in lighting).



Sign language is an important method of communication for deaf-dumb persons. As sign language is a well structured code gesture; each gesture has a meaning assigned to it. Your job is to design a sign language recognition system capable of recognizing the English alphabets only. Mention in each step (**data sensing, preprocessing, segmentation, feature extraction, learning and classification, etc.**) the strengths of the tools or techniques which you have incorporated in your design and explain why they will work. **If there are system parameters in your design, also mention how to pick the best set of values for them.**

Designing a sign language recognition system to recognize English alphabets involves a series of steps, from data collection to classification. Here's how each stage could be approached:

### 1. Data Sensing:

- **Tools/Techniques:** High-resolution cameras or depth sensors like those in Microsoft Kinect can be used to capture detailed gesture data. Infrared cameras could also be employed to reduce dependency on ambient light.
- **Strengths:** These sensors can accurately capture dynamic gestures in real-time and provide 3D spatial data, which is crucial for distinguishing subtle differences in signs.
- **Parameters:** Camera resolution and frame rate are critical. A high frame rate is necessary to capture fast gestures without blurring.

### 2. Preprocessing:

- Tools/Techniques:** Noise reduction filters, normalization techniques, and background subtraction can be used to enhance the quality of the input data.
- Strengths:** Preprocessing improves the system's robustness by reducing the impact of variations in lighting, background, and noise.
- Parameters:** The thresholds for filters and background subtraction would be tuned based on the noise levels and lighting conditions of the input data.

### 3. Segmentation:

- Tools/Techniques:** Color-based segmentation for skin detection or motion-based segmentation can be used to isolate the hands from the rest of the image.
- **Strengths:** Segmentation focuses the system on the relevant parts of the image, reducing computational complexity and improving accuracy.
- **Parameters:** Color thresholds for skin detection would be determined by the skin tone range of the users.

### 4. Feature Extraction:

- **Tools/Techniques:** Convolutional Neural Networks (CNNs) can automatically learn the most discriminative features from the hand images.
- **Strengths:** CNNs are capable of capturing spatial hierarchies in images and are invariant to translation. They can also handle rotational and scale variations if appropriately trained.
- Parameters:** The number of layers, filter sizes, and types of pooling in the CNN architecture need to be empirically determined through validation performance.



## 5. Learning and Classification:

- **Tools/Techniques:** Support Vector Machines (SVM) for a traditional approach or Recurrent Neural Networks (RNNs)/Long Short-Term Memory networks (LSTMs) for sequences.

- **Strengths:** SVMs are effective for high-dimensional data, while RNNs/LSTMs can capture temporal dependencies in sequential data, which is vital for dynamic gestures.

- **Parameters:** The kernel type in SVM or the number of hidden units and layers in RNNs/LSTMs would be optimized using grid search with cross-validation.

## 6. Validation and Testing:

- **Tools/Techniques:** Use a held-out test set or perform k-fold cross-validation to assess the model's performance.

- **Strengths:** This helps ensure that the model generalizes well to unseen data, and it reduces the likelihood of overfitting.

- **Parameters:** Cross-validation folds and evaluation metrics (accuracy, F1 score, etc.) would be chosen to provide a thorough assessment of model performance.

## 7. Post-Processing:

- **Tools/Techniques:** A smoothing algorithm or a **Hidden Markov Model** could be used to ensure temporal consistency between sequential predictions.

- **Strengths:** This step would improve the coherence of the recognized gestures, especially in continuous sign language recognition.

## 8. User Feedback Loop:

- **Tools/Techniques:** Incorporate user feedback to allow the system to learn from its mistakes and adapt to individual signing styles.

- **Strengths:** This can greatly increase the accuracy over time and personalize the system for the user's unique gesture patterns.

For picking the best set of parameter values, techniques like grid search, random search, or Bayesian optimization can be used in conjunction with cross-validation. This approach systematically tests different combinations of parameters to find the ones that yield the best validation performance.

Remember, the success of such a system also relies heavily on the quality and diversity of the data used for training. Including a wide range of signers with different hand shapes, sizes, and signing styles will create a more robust and generalizable system.

Suppose, you want to build an email spam filter using a pattern recognition technique. Here, messages are only classified according to whether they are spam or non-spam email. After learning to do this depending on a training set, we want our email reader to automatically filter out the spam messages and perhaps place them in a separate mail folder. Which pattern recognition technique do you prefer to use? Model the email filter depending on your learning technique.

Building an email spam filter is a classic binary classification problem, where the goal is to categorize emails into two classes: spam and non-spam (also known as "ham"). A variety of pattern recognition techniques could be suitable for this task, but one effective and commonly used method is the Naive Bayes classifier. This probabilistic classifier applies Bayes' theorem with strong (naive) independence assumptions between the features.

### ### Preferred Pattern Recognition Technique: Naive Bayes Classifier

#### Why Naive Bayes?

- It is simple to implement and understand.
- Despite its simplicity, it often performs well for text classification tasks such as spam filtering.
- It requires a small amount of training data to estimate the necessary parameters.
- Naive Bayes works well with high-dimensional input spaces, as is common in text classification.

### ### Modeling the Email Filter

#### 1. Data Preprocessing:

- *Tokenization*: Split the text of the emails into tokens (usually words).
- *Stop Word Removal*: Eliminate common words that don't carry much meaning.
- *Stemming/Lemmatization*: Reduce words to their base or root form.
- *Feature Extraction*: Convert text tokens into numerical features. Common approaches include the bag-of-words model and TF-IDF (Term Frequency-Inverse Document Frequency) weighting.

#### 2. Model Training:

- *Class Prior Estimation*: Calculate the prior probabilities of each class (spam and ham) based on their frequencies in the training set.
- *Conditional Probability Estimation*: Estimate the conditional probability of each token given the class using the training data.
- **Laplace Smoothing**: Apply smoothing to handle tokens that may not have occurred in the training data to avoid zero probabilities.

### **3. Classification:**

- When a new email arrives, apply the same preprocessing and feature extraction steps.
- Use the learned Naive Bayes model to calculate the posterior probability that the email is spam given its tokens.
- Classify the email as spam if the posterior probability exceeds a certain threshold, otherwise classify it as ham.

### **4. Performance Optimization:**

- Use cross-validation to estimate the performance of the classifier on unseen data.
- Adjust the threshold for classifying an email as spam or ham to balance the false positives (ham incorrectly classified as spam) and false negatives (spam incorrectly classified as ham).
- Experiment with different feature extraction techniques and model variations to improve performance.

### **5. System Deployment:**

- Integrate the trained model into the email system, so it can automatically filter incoming messages.
- Regularly update the model with new training data to adapt to changing spam tactics.

### **6. Feedback Loop:**

- Implement a feedback mechanism where users can correct misclassifications, allowing the system to learn from these corrections and improve over time.

The Naive Bayes classifier is a strong choice for a spam filter due to its effectiveness and efficiency. However, there are other machine learning models, like Support Vector Machines (SVM) and deep learning-based methods, which can also be explored and may yield better performance depending on the characteristics of the dataset and the complexity of the spam detection task.

Suppose, you are a member of a team assessing Covid-19 situation who wants to be sure about the presence or absence of Covid-19 cases among students before restarting classes. Your job is to design a COVID-19 detection system from students' health and other information. It will detect the probable COVID-19 patients from the students. Mention in each step of your design (data sensing, preprocessing, feature extraction, learning and classification, etc.) and also the strengths of the tools or techniques which you have incorporated in your design and

explain why they will work. **If there are system parameters in your design, also mention how to pick the best set of values for them.**

Designing a COVID-19 detection system based on students' health and other information to assess the presence or absence of cases involves careful planning and a multi-step approach:

### ### Data Sensing

- **Data Collection:** Gather health-related data, such as temperature readings, symptoms reported over a period (cough, fatigue, loss of taste/smell), recent travel history, contact with known cases, and vaccination status.
- **Strengths:** Collecting a comprehensive set of health indicators increases the likelihood of accurate detection. Temperature sensors, symptom tracking apps, and digital surveys can provide real-time, structured data.

### ### Preprocessing

- **Data Cleaning:** Address missing or inconsistent data through imputation or removal of incomplete records.
- **Normalization:** Scale the numerical features to a standard range if needed, which helps certain algorithms perform better.
- **Strengths:** Preprocessing creates a clean, uniform dataset, which is crucial for the accuracy of machine learning models. Normalization ensures that features with larger ranges don't dominate the decision-making process of the model.

### ### Feature Extraction

- **Feature Selection:** Use statistical techniques like correlation coefficients to identify the most relevant features, or employ algorithms like Principal Component Analysis (PCA) to reduce dimensionality while retaining the most important information.
- **Strengths:** Reducing the feature space to the most relevant variables can improve model accuracy and reduce computational complexity.

### ### Learning and Classification

- **Model Selection:** Depending on the nature of the data and the task, a variety of models could be chosen, including logistic regression for a simple binary classification, or more complex models like Random Forests or Gradient Boosting Machines for handling non-linear relationships and interactions between features.
- **Strengths:** Logistic regression offers a transparent model where the contribution of each feature is clear. Random Forest and Gradient Boosting are robust to overfitting and can model complex patterns.
- **Parameter Tuning:** Use techniques like grid search, random search, or Bayesian optimization to find the best hyperparameter values for the chosen models. Cross-validation should be used to assess the performance of different parameter settings.

### ### Validation

- **Cross-Validation:** Perform k-fold cross-validation to assess how the model will generalize to an independent dataset, which is essential for a reliable detection system.
- **Performance Metrics:** Use metrics appropriate for medical diagnosis, such as accuracy, sensitivity (true positive rate), specificity (true negative rate), and the area under the receiver operating characteristic curve (AUC-ROC).

### ### System Deployment

- **Integration:** Develop a user interface for inputting new student health data and displaying predictions. Ensure the system integrates well with existing health information systems.
- **Strengths:** A user-friendly interface encourages consistent data entry and system usage. Integration with existing systems ensures that the model has access to up-to-date and comprehensive data.

### ### Feedback and Updates

- **Feedback Loop:** Implement a system where predictions are regularly compared to actual COVID-19 test results to gather feedback on system performance.
- **Model Updates:** Regularly retrain the model with new data to adapt to changes in the virus's behavior and population immunity.

- **Strengths:** Continual learning from real-world outcomes will improve the system's accuracy and reliability over time.

For an effective COVID-19 detection system, it's critical to consider the ethical implications, privacy concerns, and the necessity for high specificity and sensitivity due to the high cost of false negatives (missed COVID cases) and false positives (unnecessary quarantine or further testing). Ensuring the system is transparent, explainable, and regularly updated with the latest public health guidance is also essential.

The following scatterplot represents a dataset consisting of two classes denoted by circles and triangles respectively. Determine whether it will be possible to classify the dataset into two classes using each of the following classifiers. Also, explain your answer with appropriate reasons.

(i) Naive Bayes (ii) Single Layer Perceptron (iii) SVM

Based on the scatterplot you've provided and the task of classifying two classes denoted by circles and triangles, here's an analysis of how each classifier might perform:

## 1. Naive Bayes:

- **Performance Expectation:** Naive Bayes might struggle with this dataset if the features are **not conditionally independent**, given the class labels, which is a fundamental assumption of Naive Bayes. If the circles and triangles are close together, and their features are **dependent**, the **overlapping regions** could lead to **misclassifications**.

- **Reasoning:** The distribution of circles and triangles suggests that there might be an underlying pattern that Naive Bayes may not capture well, especially if there's a **non-linear relationship between the features**.

## 2. Single Layer Perceptron:

- **Performance Expectation:** A single-layer perceptron may not be suitable for this dataset if the classes are **not linearly separable**, which seems to be the case. The perceptron algorithm **finds a linear decision boundary**, and from the scatterplot, it appears that a simple line may not be able to separate the circles from the triangles.

- **Reasoning:** Since the classes form a more complex boundary that doesn't appear to be a straight line, the single-layer perceptron, which only models linear boundaries, would likely fail to classify the data points correctly.

## 3. Support Vector Machine (SVM):

- **Performance Expectation:** SVM is likely to perform well on this dataset, especially if you use an **SVM with a non-linear kernel (like the Radial Basis Function (RBF) kernel)**. SVMs are capable of finding the **optimal hyperplane that maximizes the margin** between two classes, and with the right kernel, they can handle datasets where the boundary between classes is not a straight line.

- **Reasoning:** The scatterplot indicates a pattern that suggests the classes could be separated **by a curved boundary**. SVMs with non-linear kernels are designed to handle such cases, making them a good choice for this classification task.

For choosing the best set of values for the system parameters in an SVM (like the C parameter or the parameters of the kernel function), techniques such as grid search with cross-validation would be employed. This would involve training the SVM with various parameter combinations on a subset of the data (training set) and validating on another subset (validation set) to select the parameters that give the best classification performance.

In conclusion, based on the provided scatterplot, an SVM with a non-linear kernel seems to be the most appropriate choice for classifying this dataset into two classes. The Naive Bayes classifier and the single-layer perceptron may not perform as well due

to their respective limitations with respect to the apparent non-linearity and potential feature dependency in the data.